

Lecture 15: Nonparemetric regression I

Pratheepa Jeganathan

05/06/2019

Recall

- ▶ One sample sign test, Wilcoxon signed rank test, large-sample approximation, median, Hodges-Lehman estimator, distribution-free confidence interval.
- ▶ Jackknife for bias and standard error of an estimator.
- ▶ Bootstrap samples, bootstrap replicates.
- ▶ Bootstrap standard error of an estimator.
- ▶ Bootstrap percentile confidence interval.
- ▶ Hypothesis testing with the bootstrap (one-sample problem.)
- ▶ Assessing the error in bootstrap estimates.
- ▶ Example: inference on ratio of heart attack rates in the aspirin-intake group to the placebo group.
- ▶ The exhaustive bootstrap distribution.

- ▶ Discrete data problems (one-sample, two-sample proportion tests, test of homogeneity, test of independence).
- ▶ Two-sample problems (location problem - equal variance, unequal variance, exact test or Monte Carlo, large-sample approximation, H-L estimator, dispersion problem, general distribution).
- ▶ Permutation tests (permutation test for continuous data, different test statistic, accuracy of permutation tests).
- ▶ Permutation tests (discrete data problems, exchangeability.)
- ▶ Rank-based correlation analysis (Kendall and Spearman correlation coefficients.)
- ▶ Rank-based regression (straight line, multiple linear regression, statistical inference about the unknown parameters, nonparametric procedures - does not depend on the distribution of error term.)
- ▶ Smoothing (density estimation, bias-variance trade-off, curse of dimensionality)

Nonparametric regression

Introduction

- ▶ Smoothers use external functions to model the functional relationship between y and x .
- ▶ External functions: lines or low order polynomial functions.
- ▶ Nonparametric
 - ▶ lack of a specific, parametric form assumed for the regression function being estimated.
 - ▶ no strong distributional assumptions on the errors.
- ▶ We will discuss the linear smoothers.
 - ▶ estimates are linear combinations of observed data.
- ▶ Local averaging, local regression, local polynomial, kernel smoothing, penalized regression.

Nonparametric regression

- ▶ We are given n pairs of observations $(x_1, Y_1), \dots, (x_n, Y_n)$.
- ▶ Regression model
 - ▶ $Y_i = r(x_i) + \epsilon_i, i = 1, \dots, n$.
 - ▶ Y response variable.
 - ▶ x covariate/feature.
 - ▶ $\mathbb{E}(\epsilon_i) = 0$.
 - ▶ r is a regression function.
- ▶ Estimation
 - ▶ Assume covariate value x_i are fixed.
 - ▶ If we treat x_i as **random** :
 - ▶ **Data: $(X_1, Y_1), \dots, (X_n, Y_n)$.**
 - ▶ **$r(x_i) = \mathbb{E}(Y|X = x)$, mean of Y conditional on $X = x$.**
 - ▶ Assume $\mathbb{V}(\epsilon_i) = \sigma^2$ does not depend on x .
 - ▶ Estimate of $r(x)$ is $\hat{r}_n(x)$, smoother.

Linear smoother

Linear smoother

- ▶ Linear smoothers: estimates are linear combinations of observed data.
- ▶ An estimator \hat{r}_n of r is a linear smoother if, for each x , \exists a vector $\mathbf{l}(x) = (l_1(x), \dots, l_n(x))^T$ such that

$$\hat{r}_n(x) = \sum_{i=1}^n l_i(x) Y_i.$$

- ▶ Define the vector of fitted values

$$\mathbf{r} = (\hat{r}_n(x_1), \dots, \hat{r}_n(x_n))^T.$$

- ▶ It follows that

$$\mathbf{r} = \mathbf{L}\mathbf{Y},$$

where $\mathbf{Y} = (Y_1, \dots, Y_n)^T$ and $L_{ij} = l_j(x_i)$, $i = j = 1, \dots, n$ and \mathbf{L} is an $n \times n$ matrix.

Linear smoother

- ▶ The i -th row in \mathbf{L} is the weights given to each Y_i in forming the estimate $\hat{r}_n(x_i)$.
- ▶ L is called the smoothing matrix or the hat matrix.
- ▶ The i -th row of \mathbf{L} - effective kernel for estimating $r(x_i)$.
- ▶ $\nu = \text{tr}(\mathbf{L})$ - effective degrees of freedom.
- ▶ For all x , $\sum_{i=1}^n l_i(x) = 1$ (i.e., if $Y_i = c \quad \forall i$, then $\hat{r}_n(x) = c$.)

Some linear smoothers

Regressogram

- ▶ From **W2006**.
- ▶ Mostly like histogram.
- ▶ Suppose $a \leq x \leq b$, $i = 1, \dots, n$.
- ▶ Dived (a, b) in to B_1, \dots, B_m equally spaced bins.
- ▶ Let k_j be number of points in B_j . \hat{r}_n is obtained by averaging Y_i 's over each bin.

$$\hat{r}_n(x) = \frac{1}{k_j} \sum_{i: x_i \in B_j} Y_i \text{ for } x \in B_j.$$

- ▶ We can write $\hat{r}_n(x) = \sum_{i=1}^n l_i(x) Y_i$.

Regressogram

- ▶ From **W2006** Page 67, 5.24 Example.
- ▶ Example: Let $n = 9$, $m = 3$ and $k_1 = k_2 = k_3 = 3$. Then,

$$\mathbf{L} = \frac{1}{3} \times \begin{pmatrix} 1 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 1 & 1 \end{pmatrix}$$

- ▶ Effective number of freedom $\nu = \text{tr}(\mathbf{L})$.
- ▶ Binwidth $h = \frac{b - a}{m}$ controls the smoothness of the estimate.


Regressogram

- ▶ **HWC** (Page 662) Example 14.2: Nitrogen Oxide Concentrations
 - ▶ Brinkman (1981) collected data on the nitrogen oxide concentrations (Y) found in engine exhaust for ethanol engines with various equivalence ratios (x).

```
library(NSM3)
data("ethanol")
library(HoRM)
library(ggplot2)
p = regressogram(ethanol$E, ethanol$NOx,
  nbins = 20, show.bins = TRUE,
  show.means = TRUE, show.lines = TRUE,
  x.lab = "E", y.lab = "NOx",
  main = "NOx and ethanol engines metric") +
  theme(plot.title = element_text(hjust = 0.5))
```

Regressogram

NOx and ethanol engines metric



► blue dots are bins.

Local averaging (Friedman)

- ▶ **HWC** Chapter 14.1
- ▶ Estimate of r at the point x_i is taken to be the average of observed values Y_j corresponding to values x_j in some vicinity of x_i .
- ▶ The neighborhood of x_i is chosen to be the smallest symmetric window about x_i containing fixed number of observations.
- ▶ The average is a linear combination of the points in the neighborhood, thus, the fit is a linear smoother.
- ▶ `supsmu` function in R.

Local averaging

- ▶ From **W2006** Page 68, 5.26 Example.
- ▶ For $h > 0$ and let $B_x = \{i : |x_i - x| \leq h\}$.
- ▶ Let n_x be the number of points in B_x .
- ▶ $\hat{r}_n(x) = \frac{1}{n_x} \sum_{i \in B_x} Y_i$.
- ▶ We can write

$$\hat{r}_n(x) = \sum_{i=1}^n l_i(x) Y_i, \quad (1)$$

where $l_i(x) = \frac{1}{n_x}$ if $|x_i - x| \leq h$ and 0 otherwise.

- ▶ Example: Suppose $n = 9$, $x_i = \frac{i}{9}$ and $h = \frac{1}{9}$.

Local averaging




$$\mathbf{L} = \begin{pmatrix} 1/2 & 1/2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1/3 & 1/3 & 1/3 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1/3 & 1/3 & 1/3 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1/3 & 1/3 & 1/3 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1/3 & 1/3 & 1/3 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1/3 & 1/3 & 1/3 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1/3 & 1/3 & 1/3 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1/3 & 1/3 & 1/3 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1/2 & 1/2 \end{pmatrix}$$

Local averaging

```
fit.local.avg = supsmu(ethanol$E,  
  ethanol$NOx, span = "cv")  
df.fit.local.avg = data.frame(E.fit = fit.local.avg$x,  
  NOx.fit = fit.local.avg$y)  
library(dplyr)  
p = ggplot() +  
  geom_point(data = ethanol, aes(x = E, y=NOx)) +  
  geom_line(data = df.fit.local.avg,  
    aes(x = E.fit, y = NOx.fit), color = "red")
```

Local averaging



Local regression (Cleveland)

- ▶ **HWC** Chapter 14.2
- ▶ Estimate r by performing a local linear regression (locally weighted least squares) on the observations (x, Y) near x_i .
- ▶ The regression is a weighted regression - weights are related to the distance of the points used in the regression to the point x_i .
- ▶ `loess` function in R, `loess.as` for cross-validation and finding optimal span.
 - ▶ The weight function used in `loess` is tricube function:

$$W(x) = \left(1 - |x|^3\right)^3, |x| < 1.$$

- ▶ Let w_1^i, \dots, w_n^i be the weights determined by the centered and scaled W for a particular point x_i .
- ▶ The weighted local regression is found by minimizing


$$\sum_{j=1}^n w_j^i (Y_j - \beta_0^i - \beta_1^i x_j)^2.$$

- ▶ β_0^i and β_1^i are the intercept and slope of the linear relation between x and Y in the neighborhood of x_i .

Local regression

```
fit.local.lin.reg = loess(NOx ~ E, data=ethanol,  
  degree=1, span=0.19)  
  
df.fit.local.lin.reg = data.frame(E = ethanol$E,  
  NOx.fit = fit.local.lin.reg$fitted)  
library(dplyr)  
p = ggplot() +  
  geom_point(data = ethanol, aes(x = E, y=NOx)) +  
  geom_line(data = df.fit.local.lin.reg,  
    aes(x = E, y = NOx.fit), color = "red")
```

Local regression



Local polynomial regression

- ▶ **HWC** comment 8, page 665.
- ▶ Use local polynomial regression in place of linear regression.
- ▶ Let w_1^i, \dots, w_n^i be the weights determined by the centered and scaled W for a particular point x_i .
 - ▶ The weighted local polynomial regression is found by minimizing


$$\sum_{j=1}^n w_j^i (Y_i - \beta_0^i - \beta_1^i x_j - \dots - \beta_d^i x_j^d)^2.$$

- ▶ `loess` function in R allows for degrees of $d = 0, 1, 2$.

Local polynomial regression

```
fit.local.poly.reg = loess(NOx ~ E, data=ethanol,  
  degree=2, span=0.2)  
  
df.fit.local.poly.reg = data.frame(E = ethanol$E,  
  NOx.fit = fit.local.poly.reg$fitted)  
  
p = ggplot() +  
  geom_point(data = ethanol, aes(x = E, y=NOx)) +  
  geom_line(data = df.fit.local.poly.reg,  
    aes(x = E, y = NOx.fit), color = "red")
```

Local polynomial regression



Kernel smoothing

- ▶ **HWC** Chapter 14.3.
- ▶ This is not a nearest neighbor method for a given kernel K and bandwidth h .
 - ▶ The number of observations used in the estimate at any point x is not fixed but the window size is.
 - ▶ The bandwidth h is the changing value over which the least squares are minimized, rather than the span.
 - ▶ Weights are determined by how close each observation x_j to point x , bandwidth, and the kernel K .
- ▶ `npreg` from package `np`.
 - ▶ `npregbw` from package `np` for bandwidth selection.

Kernel smoothing

- ▶ **W 2006** Chapter 5.4.
- ▶ Let $h > 0$ - bandwidth.
- ▶ Nadaraya (1964, 1965) and Watson (1964):

$$\hat{r}_n(x) = \sum_{i=1}^n l_i(x) Y_i, \quad (2)$$

where K is a kernel and

$$l_i(x) = \frac{K\left(\frac{x - x_i}{h}\right)}{\sum_{j=1}^n K\left(\frac{x - x_j}{h}\right)}.$$

- ▶ The local average regression in (1) is a kernel estimator based on the boxcar kernel.
- ▶ We can show that kernel smoother is a linear smoother as in (2).

Kernel smoothing

- ▶ The choice of kernel K is not too important.
- ▶ Risk is sensitive for h_n which controls the amount of smoothing and depends on sample size n .
 - ▶ Small h_n gives rough estimates.
 - ▶ Larger h_n 's give smoother estimates.

Kernel smoothing

- ▶ An example to show the bandwidth affects the estimate.
- ▶ Let x_1, x_2, \dots, x_n be random draws from some density f .
- ▶ The risk (integrated squared error loss) of the Nadaraya-Watson kernel estimator is

$$\begin{aligned} R(\hat{r}_n, r) = & \frac{h_n^4}{4} \left(\int x^2 K(x) dx \right)^2 \int \left(r''(x) + 2r'(x) \frac{f'(x)}{f(x)} \right)^2 dx \\ & + \frac{\sigma^2 \int K^2(x) dx}{nh_n} \int \frac{1}{f(x)} dx + o(nh_n^{-1}) + o(h_n^4) \end{aligned} \quad (3)$$

as $h_n \rightarrow 0$ and $nh_n \rightarrow \infty$.



- ▶ **Design bias:** $2r'(x) \frac{f'(x)}{f(x)}$ The bias term in (3) depends on the distribution of x_i 's.

Kernel smoothing

- ▶ The optimal bandwidth will depend on the unknown function r . So we can use cross-validation to find the optimal bandwidth h^* .

Kernel smoothing

- ▶ Kernel estimators have high bias near the boundaries called **boundary bias**.



Source: Hastie, Tibshirani, Friedman (2009)

Kernel smoothing

- ▶ Alleviate the boundary bias and design bias using local polynomial regression.
 - ▶ Use the kernel K as the weight in the local polynomial regression.
 - ▶ Estimate is a linear smoother.

Kernel smoothing

```
library(np)
ethanol.npreg <- npreg(bws=.09,
  txdat=ethanol$E,
  tydat=ethanol$NOx,
  ckertype="epanechnikov")
ethanol.npreg2 <- npreg(bws=.03,
  txdat=ethanol$E,
  tydat=ethanol$NOx,
  ckertype="epanechnikov")
ethanol.npreg$MSE
```

```
## [1] 0.2816102
```


```
ethanol.npreg2$MSE
```

```
## [1] 0.1060051
```

Kernel smoothing

```
ethanol.npreg.fit = data.frame(E = ethanol$E,  
  NOx = ethanol$NOx,  
  kernel.fit = fitted(kernel.npreg),  
  kernel.fit2 = fitted(kernel.npreg2))  
  
ggplot(ethanol.npreg.fit) +  
  geom_point(aes(x = E, y = NOx)) +  
  geom_line(aes(x = E, y = kernel.fit), color = "red") +  
  geom_line(aes(x = E, y = kernel.fit2), color = "blue")
```

Kernel smoothing



Penalized regression

- ▶ **W2006** Chapter 5.5

- ▶ $Y_i = r(x_i) + \epsilon_i.$

- ▶ Suppose we estimate r by choosing $\hat{r}_n(x)$ to minimize the sum of squares

$$\sum_{i=1}^n (Y_i - \hat{r}_n(x))^2.$$

- ▶ Minimizing over all linear functions gives least squares estimator.
 - ▶ Minimizing over all functions yields a function that interpolate the data.
- ▶ To avoid the above two extreme solutions
 - ▶ locally weighted sums of squares (local averages, local linear/polynomial regression, kernel smoother).
 - ▶ minimize the penalized sums of squares.

Penalized regression

- ▶ Compute \hat{r}_n by minimizing penalized sums of squares

$$M(\lambda) = \sum_i (Y_i - \hat{r}_n(x_i))^2 + \lambda J(r),$$

where

$$J(r) = \int (r''(x))^2 dx.$$

- ▶ When $\lambda = 0$, the solution is interpolating function.
- ▶ When $\lambda \rightarrow \infty$, \hat{r}_n converges to the least squares line.
- ▶ What does \hat{r}_n looks like for $0 < \lambda < \infty$?

Splines

- ▶ A spline is a special piece-wise polynomial.
- ▶ A cubic spline
 - ▶ Let $\zeta_1, \zeta_2, \dots, \zeta_k$ be a set of ordered points - called knots - contained in some interval (a, b) .
 - ▶ A cubic spline is a continuous function r such that (i) r is a cubic polynomial over $(\zeta_1, \zeta_2), \dots$ and (ii) r has first and second derivatives at knots.

Smoothing splines

- ▶ The function $\hat{r}_n(x)$ that minimizes $M(\lambda)$ with penalty $J(r)$ is a natural cubic spline with knots at the data points.
 - ▶ \hat{r}_n does not have an explicit form.
 - ▶ Smoothing splines.
- ▶ Build an explicit basis using B-splines

$$\hat{r}_n(x) = \sum_{j=1}^N \hat{\beta}_j B_j(x),$$

- ▶ where B_1, \dots, B_N are a basis for B-splines with $N = n + 4$.
- ▶ Now we only need to find the coefficients $\hat{\beta} = (\hat{\beta}_1, \dots, \hat{\beta}_N)^T$.

B-Splines

- ▶ By expanding r in the basis we can now rewrite the minimization as follows:

$$\text{minimize } (Y - \mathbf{B}\beta)^T (Y - \mathbf{B}\beta) + \lambda \beta^T \Omega \beta,$$

where $\mathbf{B}_{ij} = B_j(X_i)$ and $\Omega_{ij} = \int B_j''(x) B_k''(x) dx$.




$$\hat{\beta} = (\mathbf{B}^T \mathbf{B} + \lambda \Omega)^{-1} \mathbf{B}^T Y.$$

- ▶ The smoothing spline is a linear smoother:

$$r = (\mathbf{B}^T \mathbf{B} + \lambda \Omega)^{-1} \mathbf{B}^T Y = \mathbf{L} Y.$$

B-Splines

- ▶ Cubic B-spline basis using nine equally spaced knots on $(0,1)$.



Source: Wasserman (2006)

Splines (Example)

```
library(splines)
```

- ▶ A Cubic Spline with 3 Knots

```
range(ethanol$E)
```

```
## [1] 0.535 1.232
```

```
cubic.spline.fit = lm(NOx ~ bs(E,  
  knots = c(.75,1,1.2)),  
  data = ethanol)
```

Splines (Example)

```
summary(cubic.spline.fit)
```

```
##
```

```
## Call:
```

```
## lm(formula = NOx ~ bs(E, knots = c(0.75, 1, 1.2)), data =
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

##	-0.6686	-0.2063	0.0214	0.1579	0.8616
----	---------	---------	--------	--------	--------

```
##
```

```
## Coefficients:
```

```
##
```

```
##      (Intercept)                Estimate Std. Error t value Pr(>|t|)
```

```
## bs(E, knots = c(0.75, 1, 1.2))1 -0.22004      0.45474    -0.485 0.629
```

```
## bs(E, knots = c(0.75, 1, 1.2))2  1.70393      0.30013     5.677 <0.001
```

```
## bs(E, knots = c(0.75, 1, 1.2))3  4.56864      0.39930    11.442 <0.001
```


```
## bs(E, knots = c(0.75, 1, 1.2))4 -0.75105      0.32717    -2.296 0.023
```

```
## bs(E, knots = c(0.75, 1, 1.2))5  0.53076      0.35427     1.499 0.135
```

Splines (Example)

```
df.cubic.spline = data.frame(x = ethanol$E,  
  y = ethanol$NOx,  
  fit = fitted(cubic.spline.fit))  
p = ggplot(data = df.cubic.spline) +  
  geom_point(aes(x = x, y = y)) +  
  geom_line(aes(x = x, y = fit),  
    color = "red") +  
  geom_vline(xintercept = c(.75,1,1.2),  
    color = "green")
```


Splines (Example)



Smoothing spline (Example)

```
smooth.spline.fit = smooth.spline(ethanol$E,  
  ethanol$NOx, cv = TRUE)  
  
df.smooth.splines = data.frame(x = smooth.spline.fit$x,  
  fit.smooth.spline = smooth.spline.fit$y)  
  
p = ggplot() +  
  geom_point(data = df.cubic.spline,  
    aes(x = x, y = y)) +  
  geom_line(data = df.cubic.spline,  
    aes(x = x, y = fit), color = "red") +  
  geom_line(data = df.smooth.splines,  
    aes(x = x, y = fit.smooth.spline),  
    color = "brown")
```

Smoothing spline (Example)



References for this lecture

HWC Chapter 14 (smoothing)

W Chapter 5