

## Lecture 27: Selection

Pratheepa Jeganathan

11/22/2019

# Recap

- ▶ What is a regression model?
- ▶ Descriptive statistics – graphical
- ▶ Descriptive statistics – numerical
- ▶ Inference about a population mean
- ▶ Difference between two population means
- ▶ Some tips on R
- ▶ Simple linear regression (covariance, correlation, estimation, geometry of least squares)
  - ▶ Inference on simple linear regression model
  - ▶ Goodness of fit of regression: analysis of variance.
  - ▶  $F$ -statistics.
  - ▶ Residuals.
  - ▶ Diagnostic plots for simple linear regression (graphical methods).

# Recap

- ▶ Multiple linear regression
  - ▶ Specifying the model.
  - ▶ Fitting the model: least squares.
  - ▶ Interpretation of the coefficients.
  - ▶ Matrix formulation of multiple linear regression
  - ▶ Inference for multiple linear regression
    - ▶  $T$ -statistics revisited.
    - ▶ More  $F$  statistics.
    - ▶ Tests involving more than one  $\beta$ .
- ▶ Diagnostics – more on graphical methods and numerical methods
  - ▶ Different types of residuals
  - ▶ Influence
  - ▶ Outlier detection
  - ▶ Multiple comparison (Bonferroni correction)
  - ▶ Residual plots:
    - ▶ partial regression (added variable) plot,
    - ▶ partial residual (residual plus component) plot.

# Recap

- ▶ Adding qualitative predictors
  - ▶ Qualitative variables as predictors to the regression model.
  - ▶ Adding interactions to the linear regression model.
  - ▶ Testing for equality of regression relationship in various subsets of a population
- ▶ ANOVA
  - ▶ All qualitative predictors.
  - ▶ One-way layout
  - ▶ Two-way layout
- ▶ Transformation
  - ▶ Achieving linearity
  - ▶ Stabilize variance
  - ▶ Weighted least squares
- ▶ Correlated Errors
  - ▶ Generalized least squares
- ▶ Bootstrapping linear regression

Selection

# Outline (Model selection)

- ▶ In a given regression situation, there are often many choices to be made.
- ▶ Recall our usual setup

$$Y_{n \times 1} = X_{n \times p} \beta_{p \times 1} + \epsilon_{n \times 1}.$$

- ▶ Any *subset*  $A \subset \{1, \dots, p\}$  yields a new regression model

$$\mathcal{M}(A) : Y_{n \times 1} = X[:, A] \beta[A] + \epsilon_{n \times 1}$$

by setting  $\beta[A^c] = 0$ .

- ▶ **Model selection** is, roughly speaking, how to choose  $A$  among the  $2^p$  possible choices.

## Election data

Here is a dataset from the book that we will use to explore different model selection approaches.

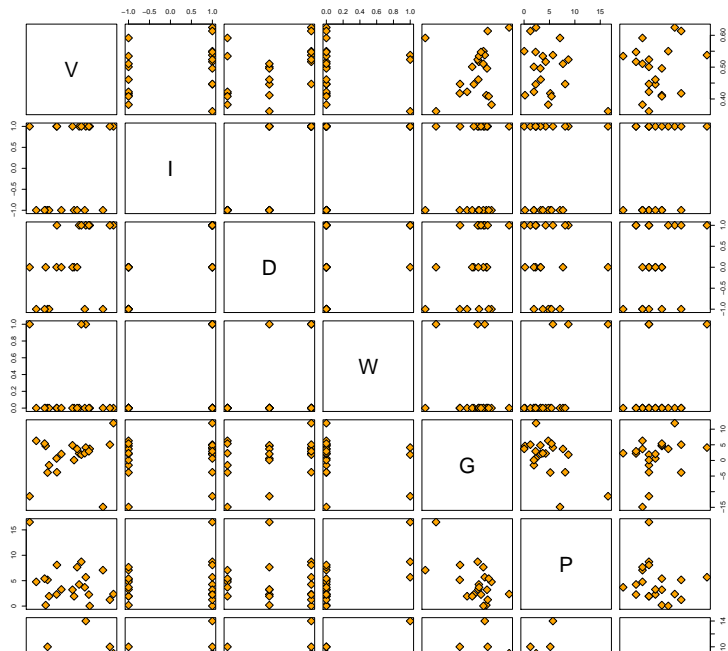
Variable	Description
<i>V</i>	votes for a presidential candidate
<i>I</i>	are they incumbent?
<i>D</i>	Democrat or Republican incumbent?
<i>W</i>	wartime election?
<i>G</i>	GDP growth rate in election year
<i>P</i>	(absolute) GDP deflator growth rate
<i>N</i>	number of quarters in which GDP growth rate $> 3.2\%$

## Election data

```
url = 'http://stats191.stanford.edu/data/election.table'  
election.table = read.table(url, header=T)  
pairs(election.table[,2:ncol(election.table)],  
      cex.labels=3, pch=23,  
      bg='orange', cex=2)
```



# Election data



# Problem & Goals

- ▶ When we have many predictors (with many possible interactions), it can be difficult to find a good model.
- ▶ Which main effects do we include?
- ▶ Which interactions do we include?
- ▶ Model selection procedures try to *simplify* / *automate* this task.
- ▶ Election data has  $2^6 = 64$  different models with just main effects!

## General comments

- ▶ This is generally an “unsolved” problem in statistics: there are no magic procedures to get you the “best model.”
- ▶ Many machine learning methods look for good “sparse” models: selecting a “sparse” model.
- ▶ “Machine learning” often work with very many predictors.
- ▶ Our model selection problem is generally at a much smaller scale than “data mining” problems.
- ▶ Still, it is a hard problem.

# Hypothetical example

- ▶ Suppose we fit a model  $F : Y_{n \times 1} = X_{n \times p} \beta_{p \times 1} + \varepsilon_{n \times 1}$  with predictors  $X_1, \dots, X_p$ .
- ▶ In reality, some of the  $\beta$ 's may be zero. Let's suppose that  $\beta_{j+1} = \dots = \beta_p = 0$ .
- ▶ Then, any model that includes  $\beta_0, \dots, \beta_j$  is *correct*: which model gives the *best* estimates of  $\beta_0, \dots, \beta_j$ ?
- ▶ Principle of *parsimony* (i.e. Occam's razor) says that the model with *only*  $X_1, \dots, X_j$  is "best".

# Justifying parsimony

- ▶ For simplicity, let's assume that  $j = 1$  so there is only one coefficient to estimate.
- ▶ Then, because each model gives an *unbiased* estimate of  $\beta_1$  we can compare models based on  $\text{Var}(\hat{\beta}_1)$ .
- ▶ The best model, in terms of this variance, is the one containing only  $X_1$ .
- ▶ What if we didn't know that only  $\beta_1$  was non-zero (which we don't know in general)?
- ▶ In this situation, we must choose a set of variables.

## Model selection: choosing a subset of variables

- ▶ To “implement” a model selection procedure, we first need a criterion or benchmark to compare two models.
- ▶ Given a criterion, we also need a search strategy.
- ▶ With a limited number of predictors, it is possible to search all possible models (leaps in R).

## Candidate criteria

# Candidate criteria

Possible criteria:

- ▶  $R^2$ : not a good criterion. Always increase with model size  
 $\implies$  “optimum” is to take the biggest model.
- ▶ Adjusted  $R^2$ : better. It “penalized” bigger models. Follows principle of parsimony / Occam’s razor.
- ▶ Mallows’s  $C_p$  – attempts to estimate a model’s predictive power, i.e. the power to predict a new observation.



## Best subsets, $R^2$

- ▶ Leaps takes a design matrix as argument: throw away the intercept column or leaps will complain.

```
election.lm = lm(V ~ I + D + W + G:I +  
                 P + N, election.table)  
#election.lm
```

Call:

```
lm(formula = V ~ I + D + W + G:I + P + N, data = election.table)
```

Coefficients:

(Intercept)	I	D	W
0.5111627	-0.0201077	0.0546159	0.0133905
P	N	I:G	
-0.0007224	-0.0051822	0.0096901	

## Best subsets, $R^2$

```
X = model.matrix(election.lm)[,-1]
library(leaps)
# Since the algorithm returns a best model of each size,
# the results do not depend on a penalty model for
# model size
# nbest: Number of subsets of each size to report
election.leaps = leaps(x = X, y = election.table$V,
  nbest=3, method='r2')
```

- Find out the predictors in the model with the largest  $R^2$ :

```
# election.leaps$which: matrix, each row can be
# used to select the columns of x in the respective model
ind = which((election.leaps$r2 == max(election.leaps$r2)))
best.model.r2 = election.leaps$which[ind, ]
best.model.r2
```

```
##      1      2      3      4      5      6
## TRUE TRUE TRUE TRUE TRUE TRUE
```

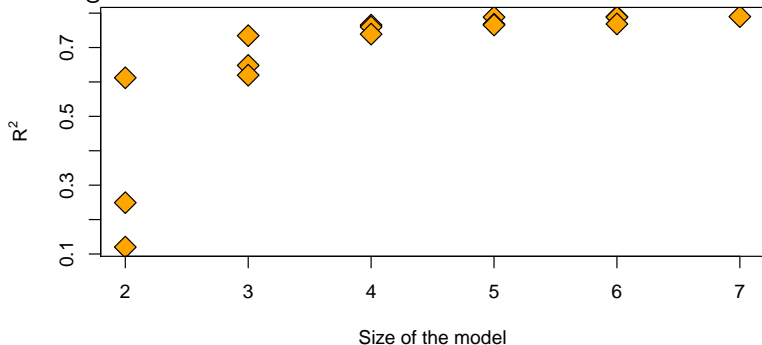
## Best subsets, $R^2$

- Let's plot the  $R^2$  as a function of the model size.

```
plot(election.leaps$size, election.leaps$r2,  
     pch=23, bg='orange', cex=2,  
     xlab = "Size of the model",  
     ylab = bquote(R2))
```

## Best subsets, $R^2$

- ▶ For example, there are three models with 2 predictors and with different  $R^2$
- ▶ We see that the full model does include all variables and has the largest  $R^2$ .



## Best subsets, adjusted $R^2$

- ▶ As we add more and more variables to the model – even random ones,  $R^2$  will increase to 1.
- ▶ Adjusted  $R^2$  tries to take this into account by replacing sums of squares by *mean squares*

$$R_a^2 = 1 - \frac{SSE/(n - p - 1)}{SST/(n - 1)} = 1 - \frac{MSE}{MST}.$$

## Best subsets, adjusted $R^2$

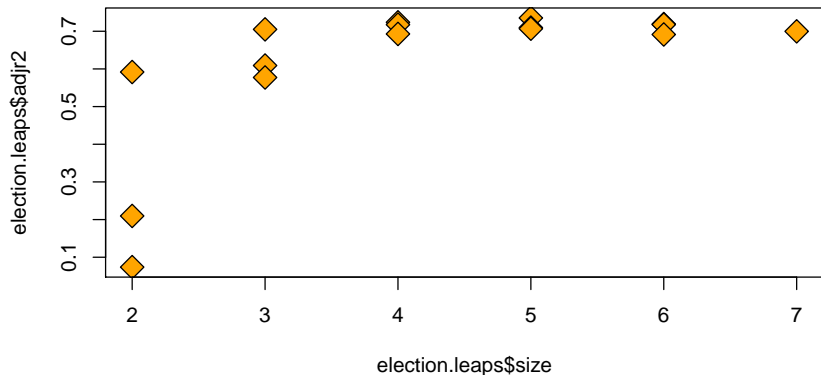
```
election.leaps = leaps(X, election.table$V, nbest=3,  
  method='adjr2')  
ind2 = which((election.leaps$adjr2 ==  
  max(election.leaps$adjr2)))  
best.model.adj2 = election.leaps$which[ind2,]  
best.model.adj2
```

##	1	2	3	4	5	6
##	TRUE	TRUE	FALSE	FALSE	TRUE	TRUE

- ▶ Best model based on the adjusted  $R^2$  has four predictor variables.

## Best subsets, adjusted $R^2$

```
plot(election.leaps$size,  
     election.leaps$adjr2,  
     pch=23, bg='orange', cex=2)
```



# Mallow's $C_p$

- ▶ Mallow's  $C_p$

$$C_p(\mathcal{M}) = \frac{SSE(\mathcal{M})}{\hat{\sigma}^2} + 2 \cdot p(\mathcal{M}) - n.$$

- ▶  $\hat{\sigma}^2 = SSE(F)/df_F$  is the “best” estimate of  $\sigma^2$  we have (use the fullest model), i.e. in the election data it uses all 6 main effects.
- ▶  $SSE(\mathcal{M})$  is the  $SSE$  of the model  $\mathcal{M}$ .
- ▶  $p(\mathcal{M})$  is the number of predictors in  $\mathcal{M}$ .
- ▶ This is an estimate of the expected mean-squared error of  $\hat{Y}(\mathcal{M})$ , it takes *bias* and *variance* of fit into account.
- ▶ Account for the sample size, effect size of the predictors, and collinearity between the predictors.



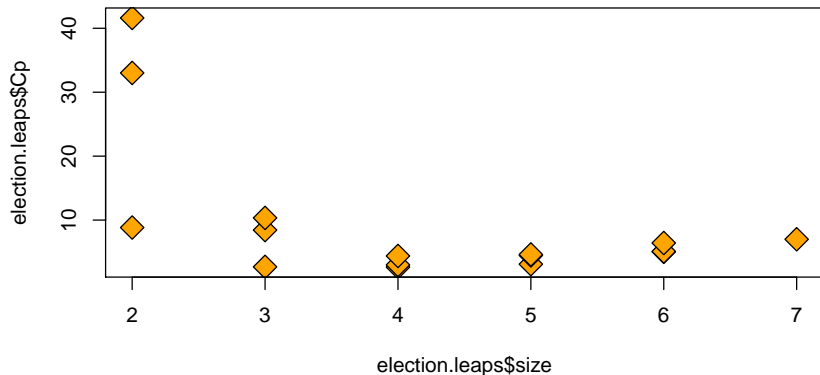
## Best subsets, Mallows's $C_p$

```
election.leaps = leaps(X, election.table$V, nbest=3,  
  method='Cp')  
indcp = which((election.leaps$Cp ==  
  min(election.leaps$Cp)))  
best.model.Cp = election.leaps$which[indcp,]  
best.model.Cp
```

##	1	2	3	4	5	6
##	FALSE	TRUE	FALSE	FALSE	TRUE	TRUE

## Best subsets, Mallows's $C_p$

```
plot(election.leaps$size,  
     election.leaps$Cp, pch=23,  
     bg='orange', cex=2)
```



## Search strategies

# Search strategies

- ▶ Given a criterion, we now have to decide how we are going to search through the possible models.
- ▶ “Best subset”: search all possible models and take the one with highest  $R_a^2$  or lowest  $C_p$  leaps. Such searches are typically feasible only up to  $p = 30$  or  $40$  at the very most.
- ▶ Stepwise (forward, backward or both): useful when the number of predictors is large. Choose an initial model and be “greedy”.
  - ▶ “Greedy” means always take the biggest jump (up or down) in your selected criterion.

# Implementations in R

- ▶ “Best subset”: use the function `leaps`. Works only for multiple linear regression models.
- ▶ Stepwise: use the function `step`. Works for any model with Akaike Information Criterion (AIC). In multiple linear regression, AIC is (almost) a linear function of  $C_p$ .

- ▶ Akaike (AIC) defined as

$$AIC(\mathcal{M}) = -2 \log L(\mathcal{M}) + 2 \cdot p(\mathcal{M})$$

where  $L(\mathcal{M})$  is the maximized likelihood of the model.

- ▶ Bayes (BIC) defined as

$$BIC(\mathcal{M}) = -2 \log L(\mathcal{M}) + \log n \cdot p(\mathcal{M})$$

- ▶ Strategy can be used for whenever we have a likelihood, so this generalizes to many statistical models.

# AIC for regression

- ▶ In linear regression with unknown  $\sigma^2$

$$-2 \log L(\mathcal{M}) = n \log(2\pi \hat{\sigma}_{MLE}^2) + n$$

where  $\hat{\sigma}_{MLE}^2 = \frac{1}{n} SSE(\hat{\beta})$

- ▶ In linear regression with known  $\sigma^2$

$$-2 \log L(\mathcal{M}) = n \log(2\pi \sigma^2) + \frac{1}{\sigma^2} SSE(\hat{\beta})$$

so AIC is very much like Mallows's  $C_p$  in this case.

# AIC for regression

- For the election data, the linear regression with all predictors has

```
n = nrow(X)
p = 7 + 1 # sigma2 is unknown
AIC_calculated = n * log(2*pi*sum(resid(election.lm)2)/n)
c(AIC_calculated, AIC(election.lm))
```

```
## [1] -66.94026 -66.94026
```



# Properties of AIC / BIC

- ▶ BIC will typically choose a model as small or smaller than AIC (if using the same search direction).
- ▶ As our sample size grows, under some assumptions, it can be shown that
  - ▶ AIC will (asymptotically) always choose a model that contains the true model, i.e. it won't leave any variables out.
  - ▶ BIC will (asymptotically) choose exactly the right model.

## Election example

- ▶ Let's take a look at step in action.
- ▶ Probably the simplest strategy is *forward stepwise* which tries to add one variable at a time, as long as it can find a resulting model whose AIC is better than its current position.
- ▶ When it can make no further additions, it terminates.

## Election example (forward stepwise)

```
# k = 2 gives the AIC, k = log(n) refers to BIC  
election.step.forward = step(lm(V ~ 1, election.table),  
  list(upper = ~ I + D + W + G + G:I + P + N),  
  direction='forward', k=2, trace=FALSE)  
election.step.forward
```

```
##
```

```
## Call:
```

```
## lm(formula = V ~ D + P, data = election.table)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)                D                P
```

```
##      0.514022      0.043134     -0.006017
```

- Summary of the chosen model based on forward stepwise and AIC.

```
##summary(election.step.forward)
```

Call:

```
lm(formula = V ~ D + P, data = election.table)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.101121	-0.036838	-0.006987	0.019029	0.163250

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.514022	0.022793	22.552	1.2e-14 ***
D	0.043134	0.017381	2.482	0.0232 *
P	-0.006017	0.003891	-1.546	0.1394

---

Signif. codes:

0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.06442 on 18 degrees of freedom

Multiple R-squared: 0.3372, Adjusted R-squared: 0.2636

F-statistic: 4.579 on 2 and 18 DF, p-value: 0.02468

# Interactions and hierarchy

- ▶ We notice that although the *full* model we gave it had the interaction  $I:G$ , the function step never tried to use it.
- ▶ This is due to some rules implemented in step that do not include an interaction unless both main effects are already in the model.
- ▶ In this case, because neither  $I$  nor  $G$  were added, the interaction was never considered.
- ▶ In the `leaps` example, we gave the function the design matrix and it did not have to consider interactions: they were already encoded in the design matrix.

## BIC example

- ▶ The only difference between AIC and BIC is the price paid per variable. This is the argument  $k$  to `step`.
- ▶ By default  $k=2$  and for BIC we set  $k=\log(n)$ .
- ▶ If we set  $k=0$  it will always add variables.

```
election.step.forward.BIC = step(lm(V ~ 1,  
election.table),  
list(upper = ~ I + D + W + G:I + P + N),  
direction='forward', k=log(nrow(X)))
```

```
## Start: AIC=-106.73
```

```
## V ~ 1
```

```
##
```

```
##           Df Sum of Sq      RSS      AIC
```

```
## + D         1 0.0280805 0.084616 -109.71
```

```
## <none>                0.112696 -106.73
```

```
## + I         1 0.0135288 0.099167 -106.38
```

```
## + P         1 0.0124463 0.100250 -106.15
```

```
## + N         1 0.0004043 0.112271 -104.15
```

# BIC example

Start: AIC=-106.73

V ~ 1

	Df	Sum of Sq	RSS	AIC
+ D	1	0.0280805	0.084616	-109.71
<none>			0.112696	-106.73
+ I	1	0.0135288	0.099167	-106.38
+ P	1	0.0124463	0.100250	-106.15
+ N	1	0.0024246	0.110271	-104.15
+ W	1	0.0009518	0.111744	-103.87

Step: AIC=-109.71

V ~ D

	Df	Sum of Sq	RSS	AIC
<none>			0.084616	-109.71
+ P	1	0.0099223	0.074693	-109.28
+ W	1	0.0068141	0.077801	-108.43
+ I	1	0.0012874	0.083328	-106.99
+ N	1	0.0000033	0.084612	-106.67

# BIC example

```
#summary(election.step.forward.BIC)
```

Call:

```
lm(formula = V ~ D, data = election.table)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.125196	-0.033002	-0.007789	0.018511	0.150298

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.48640	0.01466	33.172	<2e-16 ***
D	0.04509	0.01796	2.511	0.0212 *

---

Signif. codes:

0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.06673 on 19 degrees of freedom

Multiple R-squared: 0.2492, Adjusted R-squared: 0.2097

F-statistic: 6.305 on 1 and 19 DF, p-value: 0.02124



# Backward selection

- ▶ Let's consider backwards stepwise. This starts at a full model and tries to delete variables.
- ▶ There is also a `direction="both"` option.

```
election.step.backward = step(election.lm,  
  direction='backward')
```

```
## Start:  AIC=-128.54  
## V ~ I + D + W + G:I + P + N  
##  
##           Df Sum of Sq      RSS      AIC  
## - P         1  0.000055 0.023741 -130.49  
## - W         1  0.000170 0.023855 -130.39  
## <none>                        0.023686 -128.54  
## - N         1  0.003133 0.026818 -127.93  
## - D         1  0.011926 0.035612 -121.97  
## - I:G       1  0.050640 0.074325 -106.52  
##  
## Step:    AIC= 130.49
```

# Backward selection

Start: AIC=-128.54

V ~ I + D + W + G:I + P + N

	Df	Sum of Sq	RSS	AIC
- P	1	0.000055	0.023741	-130.49
- W	1	0.000170	0.023855	-130.39
<none>			0.023686	-128.54
- N	1	0.003133	0.026818	-127.93
- D	1	0.011926	0.035612	-121.97
- I:G	1	0.050640	0.074325	-106.52

Step: AIC=-130.49

V ~ I + D + W + N + I:G

	Df	Sum of Sq	RSS	AIC
- W	1	0.000120	0.023860	-132.38
<none>			0.023741	-130.49
- N	1	0.003281	0.027021	-129.77
- D	1	0.013983	0.037724	-122.76
- I:G	1	0.053507	0.077248	-107.71

Step: AIC=-132.38

V ~ I + D + N + I:G

	Df	Sum of Sq	RSS	AIC
<none>			0.023860	-132.38
- N	1	0.003199	0.027059	-131.74
- D	1	0.013867	0.037727	-124.76
- I:G	1	0.059452	0.083312	-108.12

# Backward selection

```
# summary(election.step.backward)
```

Call:

```
lm(formula = V ~ I + D + N + I:G, data = election.table)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.043509	-0.019208	-0.004912	0.009626	0.090627

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.506530	0.020689	24.483	4.15e-14 ***
I	-0.019417	0.014701	-1.321	0.20515
D	0.055436	0.018180	3.049	0.00765 **
N	-0.004653	0.003177	-1.465	0.16241
I:G	0.009588	0.001519	6.314	1.03e-05 ***

---

Signif. codes:

0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.03862 on 16 degrees of freedom

Multiple R-squared: 0.7883, Adjusted R-squared: 0.7353

F-statistic: 14.89 on 4 and 16 DF, p-value: 2.95e-05

# Cross-validation

- ▶ Yet another model selection criterion is  $K$ -fold cross-validation.
- ▶ Fix a model  $\mathcal{M}$ . Break data set into  $K$  approximately equal sized groups  $(G_1, \dots, G_K)$ .
- ▶ For  $(i \text{ in } 1:K)$  Use all groups except  $G_i$  to fit model, predict outcome in group  $G_i$  based on this model  $\hat{Y}_{j,\mathcal{M},G_i}, j \in G_i$ .
- ▶ Similar to what we saw in Cook's distance / DFFITS.
- ▶ Estimate  $CV(\mathcal{M}) = \frac{1}{n} \sum_{i=1}^K \sum_{j \in G_i} (Y_j - \hat{Y}_{j,\mathcal{M},G_i})^2$ .

## Comments about cross-validation.

- ▶ It is a general principle that can be used in other situations to “choose parameters.”
- ▶ Pros (partial list): “objective” measure of a model’s predictive power.
- ▶ Cons (partial list): all we know about inference is *usually* “out the window” (also true for other model selection procedures).
- ▶ If goal is not really inference about certain specific parameters, it is a reasonable way to compare models.

## Example (Cross-validation)

```
library(boot)
#Fitting Generalized Linear Models
election.glm = glm(V ~ ., data=election.table)
# 5-fold cross-validation
# The first component is the raw cross-validation
# estimate of prediction error.
# The second component is the adjusted cross-validation
# estimate.
# The adjustment is designed to compensate for
# the bias introduced by not using
# leave-one-out cross-validation.
cv.glm(model.frame(election.glm),
        election.glm, K=5)$delta

## [1] 0.01411831 0.01242346
```

## $C_p$ versus 5-fold cross-validation

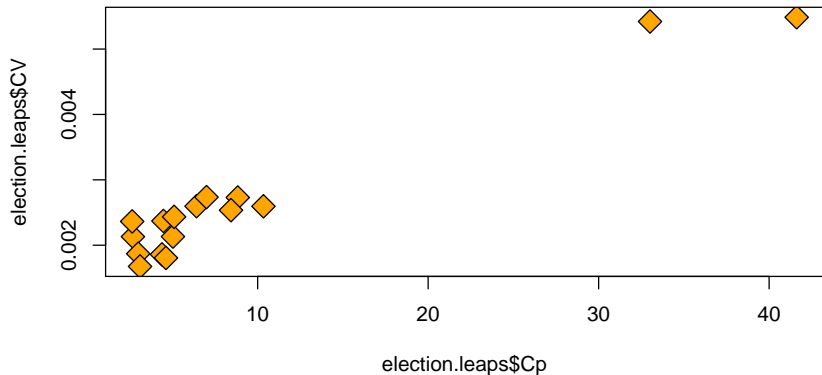
- ▶ Let's plot our  $C_p$  versus the  $CV$  score.
- ▶ Keep in mind that there is additional randomness in the  $CV$  score due to the random assignments to groups.





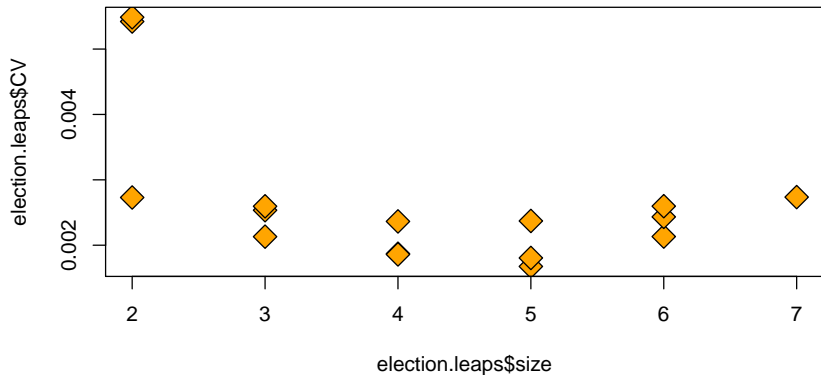
## $C_p$ versus 5-fold cross-validation

```
plot(election.leaps$Cp, election.leaps$CV,  
     pch=23, bg='orange', cex=2)
```



## $C_p$ versus 5-fold cross-validation

```
plot(election.leaps$size, election.leaps$CV,  
     pch=23, bg='orange', cex=2)
```



## $C_p$ versus 5-fold cross-validation

```
indcp_5fold = which((election.leaps$CV==  
                    min(election.leaps$CV)))  
best.model.Cv = election.leaps$which[indcp_5fold,]  
best.model.Cv
```

##	1	2	3	4	5	6
##	TRUE	TRUE	FALSE	FALSE	TRUE	TRUE

# Summarizing results

- ▶ The model selected depends on the criterion used.

Criterion	Model
$R^2$	$\sim I + D + W + G : I + P + N$
$R_a^2$	$\sim I + D + P + N$
$C_p$	$\sim D + P + N$
AIC forward	$\sim D + P$
BIC forward	$\sim D$
AIC backward	$\sim I + D + N + I : G$
5-fold CV	$\sim I + W$

- ▶ **The selected model is random and depends on which method we use!**

## Where we are so far

- ▶ Many other “criteria” have been proposed.
- ▶ Some work well for some types of data, others for different data.
- ▶ Check diagnostics!
- ▶ These criteria (except cross-validation) are not “direct measures” of predictive power, though Mallow’s  $C_p$  is a step in this direction.
- ▶  $C_p$  measures the quality of a model based on both *bias* and *variance* of the model. Why is this important?
- ▶ *Bias-variance* tradeoff is ubiquitous in statistics. More soon.

## A larger example

- ▶ Resistance of  $n = 633$  different HIV+ viruses to drug 3TC.
- ▶ Features  $p = 91$  are mutations in a part of the HIV virus, response is log fold change in vitro.

## Example (HIV and mutations)

```
X_HIV = read.table('http://stats191.stanford.edu/data/NRTI_
Y_HIV = read.table('http://stats191.stanford.edu/data/NRTI_
set.seed(0)
Y_HIV = as.matrix(Y_HIV)[,1]
X_HIV = as.matrix(X_HIV)
nrow(X_HIV)

## [1] 633
```

# Forward stepwise

```
D = data.frame(X_HIV, Y_HIV)
M = lm(Y_HIV ~ ., data=D)
M_forward = step(lm(Y_HIV ~ 1, data=D), list(upper=M),
  trace=FALSE, direction='forward')
#M_forward
```

Call:

```
lm(formula = Y_HIV ~ V68 + V17 + V19 + V23 + V54 + V67 + V82 +
  V32 + V81 + V87 + V57 + V41 + V31 + V29 + V30 + V70 + V39 +
  V26 + V69 + V40 + V62 + V64 + V80, data = D)
```

Coefficients:

(Intercept)	V68	V17	V19
0.3447	4.4731	1.5777	0.3233
	V23	V67	V82
1.4172	0.4990	0.3796	0.3854
	V32	V87	V57
0.6446	0.4113	0.5646	0.1970
	V41	V31	V29
0.5896	-0.2111	0.5407	0.6294
	V70	V39	V26
-0.1591	0.4797	-0.1633	0.2094
	V40	V62	V64
-0.3003	-0.3095	0.1792	-0.1119



# Backward stepwise

```
M_backward = step(M, list(lower= ~ 1),  
  trace=FALSE, direction='backward')  
#M_backward
```

Call:

```
lm(formula = Y_HIV ~ V17 + V19 + V23 + V26 + V29 + V30 + V31 +  
  V32 + V39 + V40 + V41 + V54 + V57 + V62 + V64 + V67 + V68 +  
  V69 + V70 + V80 + V81 + V82 + V87, data = D)
```

Coefficients:

(Intercept)	V17	V19	V23
0.3447	1.5777	0.3233	1.4172
	V26	V29	V30
			V31
-0.1633	0.5407	0.6294	-0.2111
	V32	V39	V40
			V41
0.6446	0.4797	-0.3003	0.5896
	V54	V57	V62
			V64
0.4990	0.1970	-0.3095	0.1792
	V67	V68	V69
			V70
0.3796	4.4731	0.2094	-0.1591
	V80	V81	V82
			V87
-0.1119	0.4113	0.3854	0.5646

# Both directions

```
M_both1 = step(M, list(lower= ~ 1, upper=M),  
  trace=FALSE, direction='both')  
#M_both1
```

Call:

```
lm(formula = Y_HIV ~ V17 + V19 + V23 + V26 + V29 + V30 + V31 +  
  V32 + V39 + V40 + V41 + V54 + V57 + V62 + V64 + V67 + V68 +  
  V69 + V70 + V80 + V81 + V82 + V87, data = D)
```

Coefficients:

(Intercept)	V17	V19	V23
0.3447	1.5777	0.3233	1.4172
V26	V29	V30	V31
-0.1633	0.5407	0.6294	-0.2111
V32	V39	V40	V41
0.6446	0.4797	-0.3003	0.5896
V54	V57	V62	V64
0.4990	0.1970	-0.3095	0.1792
V67	V68	V69	V70
0.3796	4.4731	0.2094	-0.1591
V80	V81	V82	V87
-0.1119	0.4113	0.3854	0.5646

# Both directions

```
M_both2 = step(lm(Y_HIV ~ 1, data=D),  
  list(lower= ~ 1, upper=M),  
  trace=FALSE, direction='both')  
#M_both2
```

Call:

```
lm(formula = Y_HIV ~ V68 + V17 + V19 + V23 + V54 + V67 + V82 +  
  V32 + V81 + V87 + V57 + V41 + V31 + V29 + V30 + V70 + V39 +  
  V26 + V69 + V40 + V62 + V64 + V80, data = D)
```

Coefficients:

(Intercept)	V68	V17	V19
0.3447	4.4731	1.5777	0.3233
V23	V54	V67	V82
1.4172	0.4990	0.3796	0.3854
V32	V81	V87	V57
0.6446	0.4113	0.5646	0.1970
V41	V31	V29	V30
0.5896	-0.2111	0.5407	0.6294
V70	V39	V26	V69
-0.1591	0.4797	-0.1633	0.2094
V40	V62	V64	V80
-0.3003	-0.3095	0.1792	-0.1119

# Compare selected models

```
sort(names(coef(M_forward)))  
sort(names(coef(M_backward)))  
sort(names(coef(M_both1)))  
sort(names(coef(M_both2)))
```

[1]	"(Intercept)"	"V17"	"V19"	"V23"
[5]	"V26"	"V29"	"V30"	"V31"
[9]	"V32"	"V39"	"V40"	"V41"
[13]	"V54"	"V57"	"V62"	"V64"
[17]	"V67"	"V68"	"V69"	"V70"
[21]	"V80"	"V81"	"V82"	"V87"
[1]	"(Intercept)"	"V17"	"V19"	"V23"
[5]	"V26"	"V29"	"V30"	"V31"
[9]	"V32"	"V39"	"V40"	"V41"
[13]	"V54"	"V57"	"V62"	"V64"
[17]	"V67"	"V68"	"V69"	"V70"
[21]	"V80"	"V81"	"V82"	"V87"
[1]	"(Intercept)"	"V17"	"V19"	"V23"
[5]	"V26"	"V29"	"V30"	"V31"
[9]	"V32"	"V39"	"V40"	"V41"
[13]	"V54"	"V57"	"V62"	"V64"
[17]	"V67"	"V68"	"V69"	"V70"
[21]	"V80"	"V81"	"V82"	"V87"
[1]	"(Intercept)"	"V17"	"V19"	"V23"
[5]	"V26"	"V29"	"V30"	"V31"
[9]	"V32"	"V39"	"V40"	"V41"
[13]	"V54"	"V57"	"V62"	"V64"
[17]	"V67"	"V68"	"V69"	"V70"
[21]	"V80"	"V81"	"V82"	"V87"

## BIC vs AIC

```
M_backward_BIC = step(M, list(lower= ~ 1), trace=FALSE,  
  direction='backward', k=log(633))  
M_forward_BIC = step(lm(Y_HIV ~ 1, data=D), list(upper=M),  
  trace=FALSE, direction='forward', k=log(633))  
M_both1_BIC = step(M, list(upper=M, lower=~1),  
  trace=FALSE, direction='both', k=log(633))  
M_both2_BIC = step(lm(Y_HIV ~ 1, data=D), list(upper=M, lower=1),  
  trace=FALSE, direction='both', k=log(633))
```

# BIC vs AIC

```
sort(names(coef(M_backward_BIC)))  
sort(names(coef(M_forward_BIC)))  
sort(names(coef(M_both1_BIC)))  
sort(names(coef(M_both2_BIC)))
```

```
[1] "(Intercept)" "V17"      "V19"      "V23"  
[5] "V29"         "V30"      "V31"      "V32"  
[9] "V41"         "V57"      "V67"      "V68"  
[13] "V81"         "V82"      "V87"      "  
[1] "(Intercept)" "V17"      "V19"      "V23"  
[5] "V31"         "V32"      "V41"      "V54"  
[9] "V57"         "V67"      "V68"      "V81"  
[13] "V82"         "V87"      "  
[1] "(Intercept)" "V17"      "V19"      "V23"  
[5] "V29"         "V30"      "V31"      "V32"  
[9] "V41"         "V57"      "V67"      "V68"  
[13] "V81"         "V82"      "V87"      "  
[1] "(Intercept)" "V17"      "V19"      "V23"  
[5] "V31"         "V32"      "V41"      "V54"  
[9] "V57"         "V67"      "V68"      "V81"  
[13] "V82"         "V87"      "
```

Inference after selection

# Inference after selection: data snooping and splitting

- ▶ Each of the above criteria return a model. The summary provides  $p$ -values.

```
summary(election.step.forward)
```

```
##
```

```
## Call:
```

```
## lm(formula = V ~ D + P, data = election.table)
```

```
##
```

```
## Residuals:
```

```
##          Min          1Q          Median          3Q          Max
```

```
## -0.101121 -0.036838 -0.006987  0.019029  0.163250
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept)  0.514022   0.022793  22.552 1.2e-14 ***
```

```
## D           0.043134   0.017381   2.482  0.0232 *
```

```
## P          -0.006017   0.003891  -1.546  0.1394
```

```
##
```



# Inference after selection

- ▶ We can also form confidence intervals. **But, can we trust these intervals or tests? No!**
- ▶ Recommended reading [Work by Jonathan Taylor](#)

```
library(selectiveInference)
```

# Reference

- ▶ **CH** Chapter 11 (Variable selection procedures)
- ▶ Lecture notes of [Jonathan Taylor](#) .